
Building a Multiplayer Online Game (MOG) using Named Data Networking (NDN)


MAI Research Project
Stefano Lupo
Supervisor: Dr Stefan Weber

Introduction

IP and NDN

IP: The “*narrow waist*” of the Internet

- IP is the *universal network layer*
- All internet connected devices speak IP
- Entirely host centric abstraction
 - Designed for point-to-point communication
 - Influenced by telephone networks of 20th Century



Named Data Networking

- Instance of Information Centric Networking (ICN)
- Move to an entirely data centric abstraction

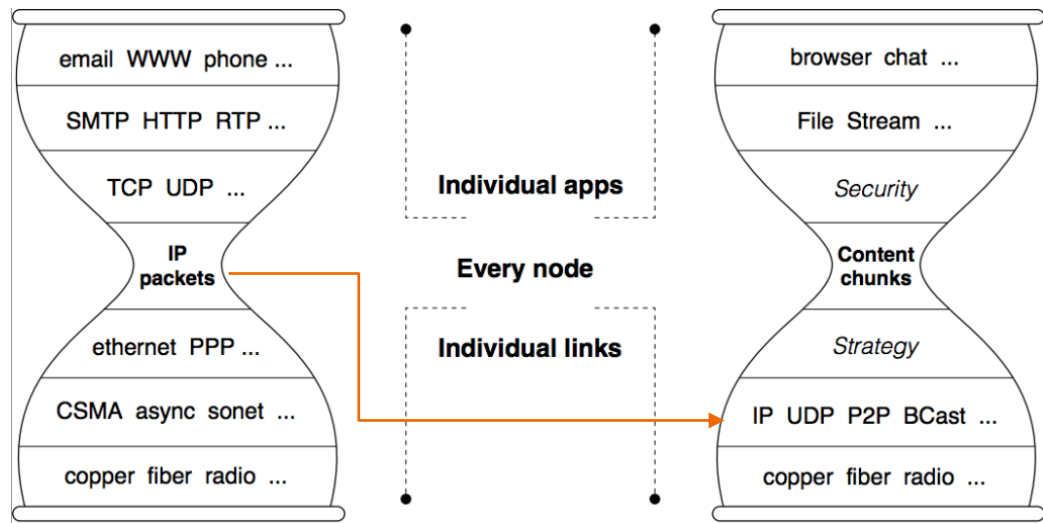


Figure: NDN Executive Summary, <https://named-data.net/project/execsummary/>

A brief overview of the State of the Art



How does NDN work?

NDN Primitives

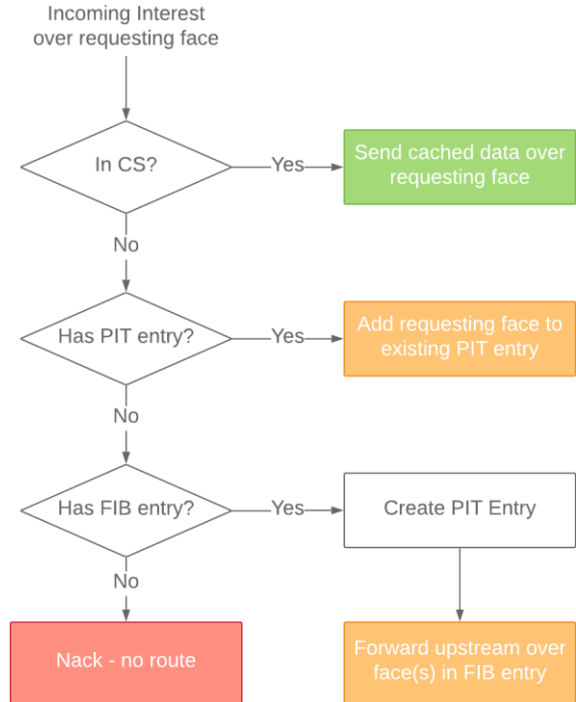
- Interest packet
 - Contains the name of data they refer to
 - Consumers request data from the network by *expressing Interests*
- Data packet
 - Produced in response to Interest packets
 - Contain the name of the data they encapsulate
 - Signed by producers
- Face abstraction
 - Physical interface - e.g. network card
 - Logical interface - e.g. application producing data under namespace

How does NDN work?

Three Key Data Structures

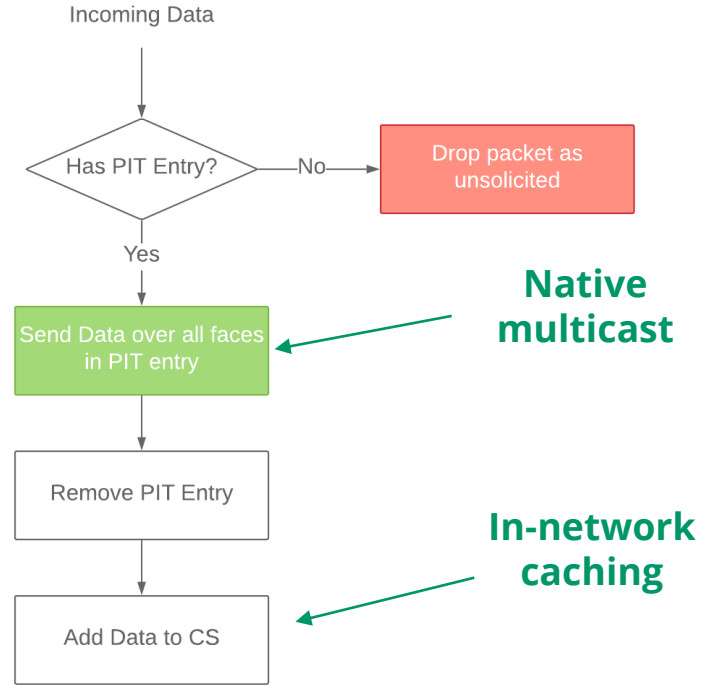
- Content Store (CS)
 - Cache of {Interests → Data} which has been recently obtained
- Pending Interests Table (PIT)
 - Stores outstanding Interests which have been forwarded upstream
 - Contains Name of Interest and list of downstream requesting faces
- Forwarding Information Base (FIB)
 - Similar to IP counterpart
 - Contains next-hop information for Interests

How does NDN work?



On Interest Received

Interest aggregation



On Data Received

Native multicast

In-network caching

Key Differences Between NDN and IP

NDN

- Packets routed by name only
- Stateful routing (FIB) and forwarding (PIT)
- Multipath forwarding
- Secures data at production time
- In-network caching

IP

- Packets routed by destination IP
- Stateful routing only
- Spanning tree forwarding
- Secures communication channel (TLS)
- No in-network caching

NDN Projects

Core software

- NDN Forwarding Daemon (**NFD***)
- The NDN Common Client Libraries (e.g. **jNDN***, pyNDN2, ndn-cxx)

Dataset synchronization

- CCNx Sync
- **ChronoSync***
- RoundSync, pSync

Routing protocol

- NDN Link State Routing Protocol (**NLSR***)

Real time applications

- Voice over CCN (VoCCN)
- Video Conferencing (NDN RTC)

MOG Architectures

Client Server (C/S)

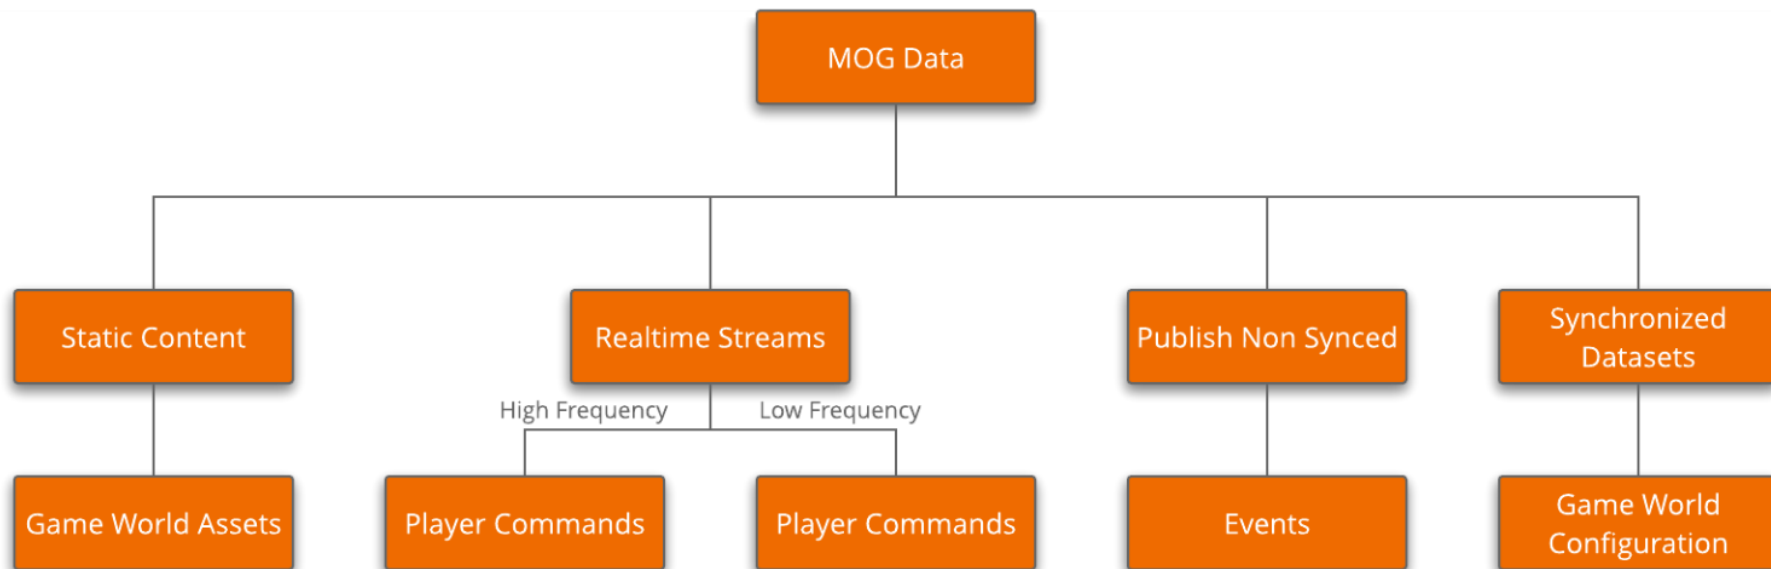
- One, or small number of, centralized servers
- Authoritative source of game state (handles all updates)
- Must send updates to all interested players (expensive)

Application Level P2P

- Distribute workload amongst players
- Peers inform each other of state changes
- Scalable, hard to maintain consistency, easier to cheat.

Game state must *appear* consistent to all connected players

MOG Data Taxonomy




Dead Reckoning and Client Side Extrapolation

- Typical frame rates are 30 or 60 FPS
- Moving remote players based on updates alone is not feasible!
- Remote updates include velocities (and sometimes accelerations!)
- Clients update their copies of remote objects
 - Use dynamics of the remote object to guess at their position between updates
- Dead reckoning convergence mechanism
 - Smoothly transition between extrapolated position and updated position
 - Must be done while continuing to extrapolate to future positions

Area of Interest Management

- Players only interested in sub region of game world
- Simple approach
 - Distance based
- More complex
 - Take game world into account
 - Precompute interest tiles for every tile a player can visit



Tile based interest filtering

Benefits of NDN in a MOG Context

- Interest aggregation at intermediate routers
 - Especially useful in P2P architecture
 - Requires n^2 connections in IP
 - Number of connections dependent on topology in NDN
- Native multicast
 - Publishers (game players) need only produce a piece of data once
 - Data is disseminated by NDN routers
- In-network caching
 - Peers can benefit from common data sent as a result of another peer's interest

Closely Related Projects

- Egal Car: A Peer-to-Peer Car Racing Game, Zening Qu, Jeff Burke
 - Uses CCNx Sync to synchronize car positions
 - No support for interaction between game objects
- Matryoshka: Design of NDN Multiplayer Online Game, Zhehao Wang, Zening Qu and Jeff Burke
 - Very limited information (only 2 page paper)
 - No testing or evaluation
- NDNGame: A NDN-based Architecture for Online Games, Diego Barros, Marcial Fernandez
 - Vague outline of general design
 - No actual implementation, no results of simulations outlined

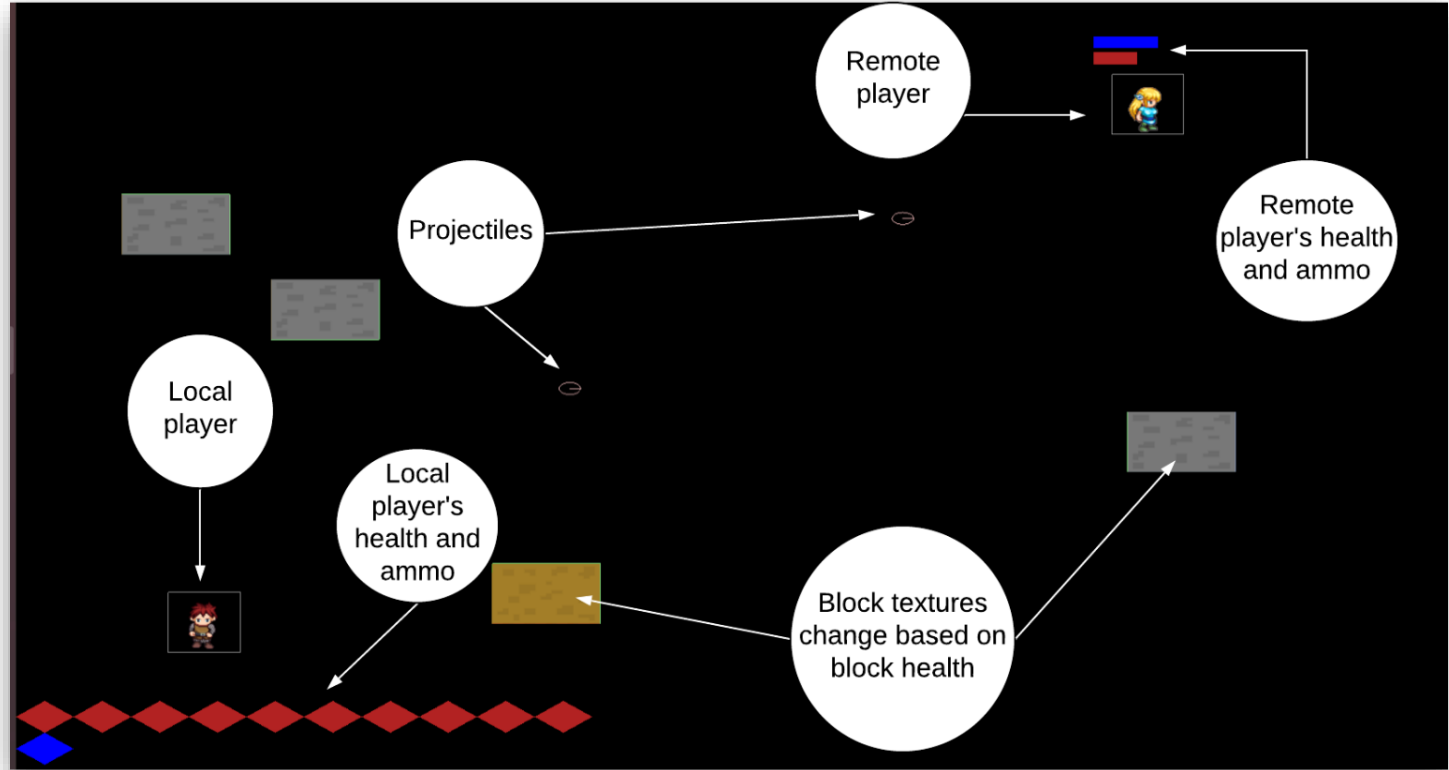
Research Questions

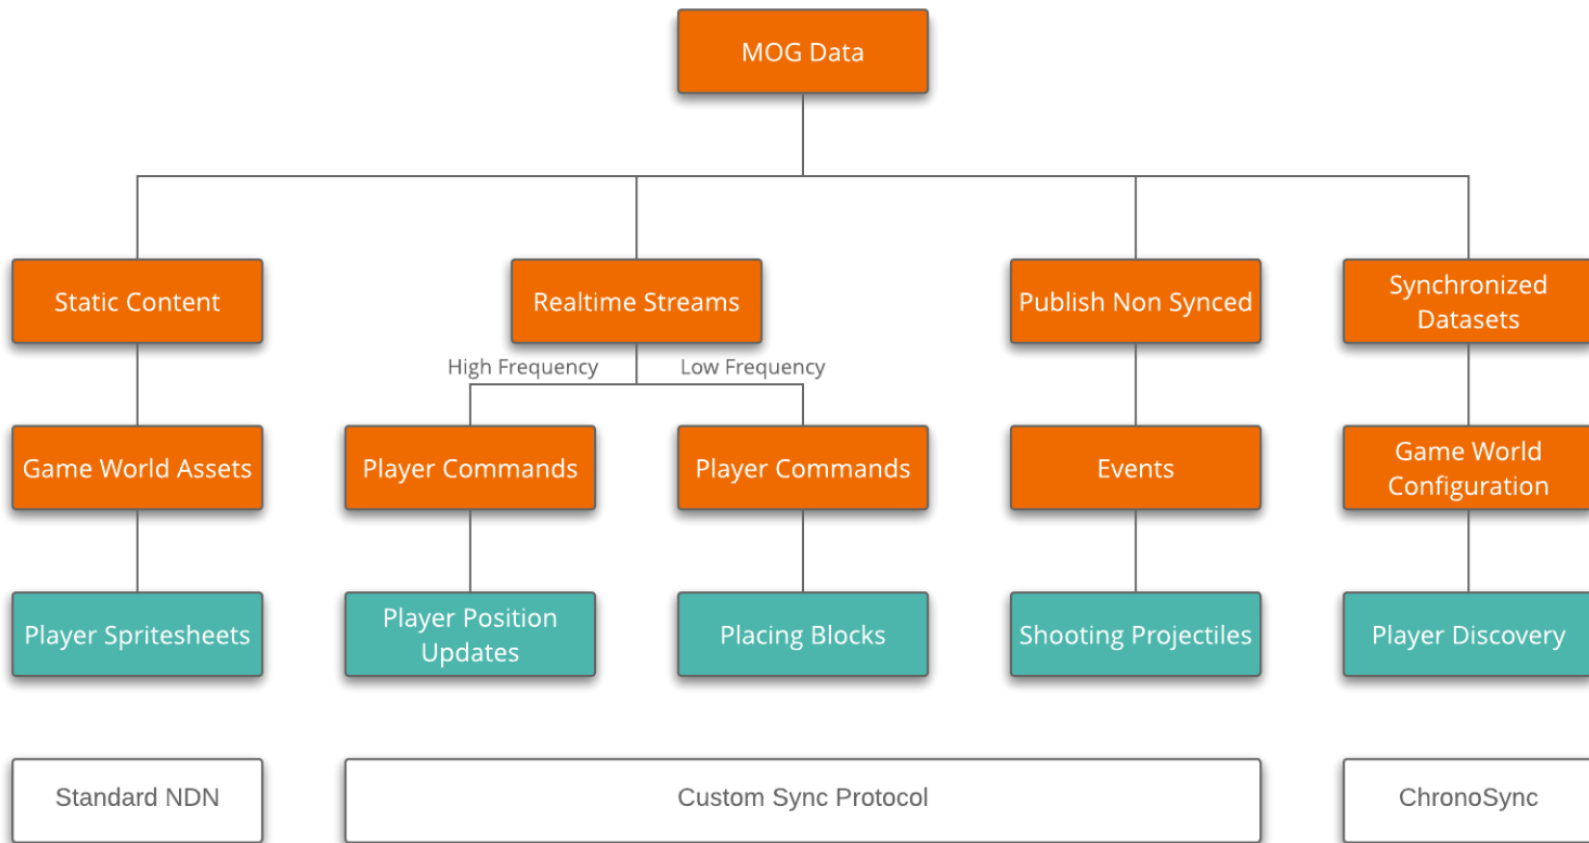
How can the benefits of NDN be exploited in a MOG context?

How can the different data types in MOGs be synced over NDN?

Design

A Basic 2D Game





Custom Sync Protocol: Motivation

Limitations of existing protocols

- Require **two round trips** (one for sync data, one for actual data)
- Write access to dataset is synchronous
 - Concurrent writes cause major problems and require reconciliation

Key Insights




- MOG data is not inherently sequential
- Publishers need only publish the **newest** instance of their data

Custom Sync Protocol: Sync

- Interest Expression Name for Sync
 - `<game_prefix>/<game_id>/<player_name>/<type>/sync/<sn>`
 - `/ndngame/0/player1/status/sync/1`
- Data Reply Name
 - `<game_prefix>/<game_id>/<player_name>/<type>/sync/<sn>/<next_sn>`
 - `/ndngame/0/player1/status/sync/1/50`
- Key points
 - Publisher responds with **Data[49]** not Data[1]
 - Subscribers obtains next sequence number and will express for Data[50] next
 - Natural catch up mechanism
 - Exploits CanBePrefix field of NDN Interest packets

Custom Sync Protocol: Benefits

- Publishers only produce Data when there are **updates**
- Interest are **aggregated** at intermediate routers
- Data produced is **multicast** back to subscribers

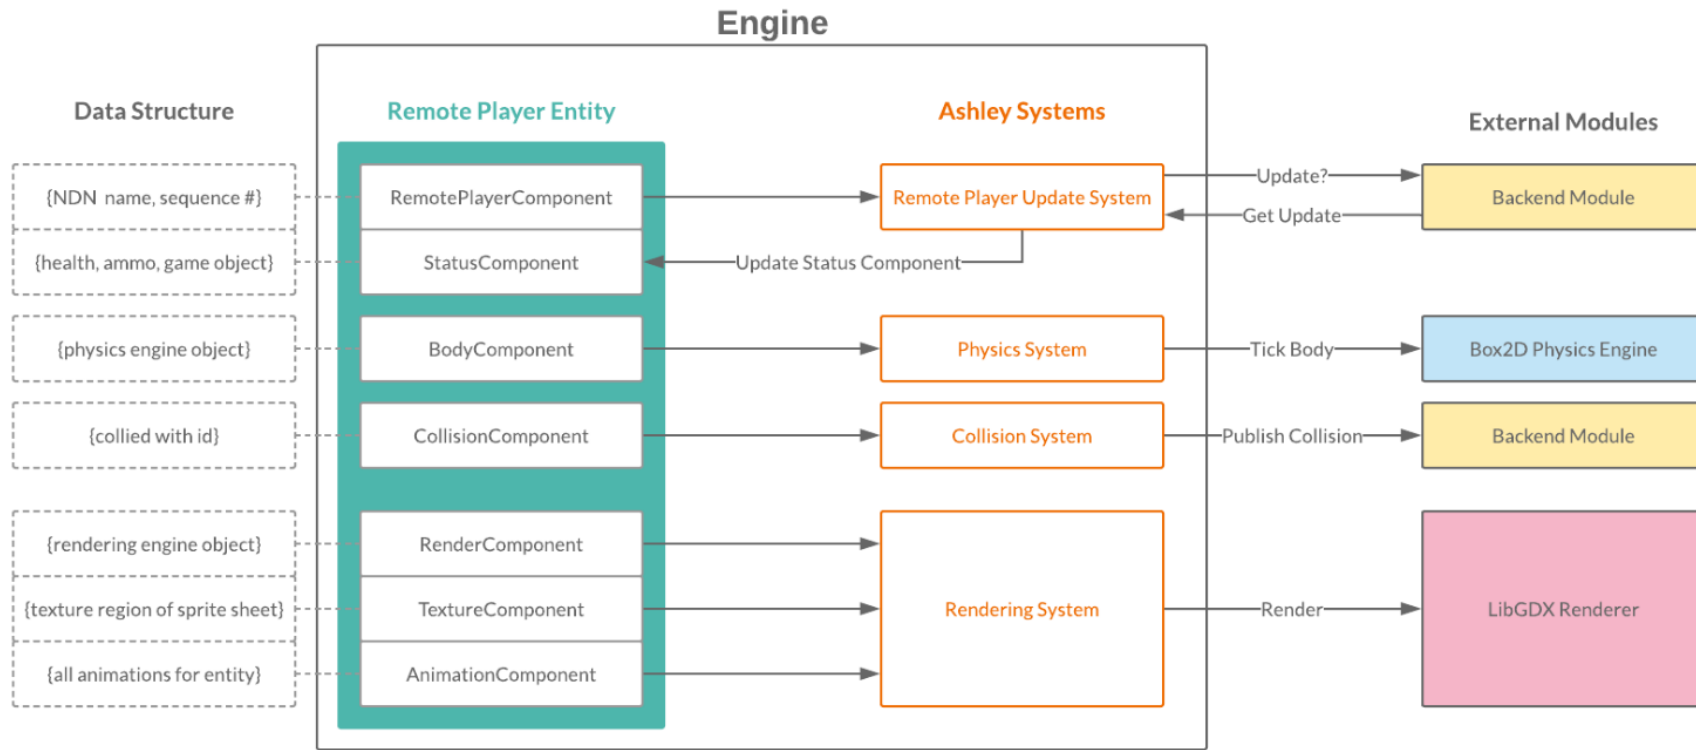


Implementation

Front End Implementation

- Built in Java using LibGDX
 - Java wrapper for OpenGL
 - Produces cross platform executables
 - Supports *headless* mode
- Uses Entity Management System (EMS) called Ashley
 - Components: Basic data structures
 - Entities: Contain a list of components
 - Systems: Performs operations on entities which contain certain components
 - Engines: Stores entities, invokes systems

Entity Representation for Remote Player



Backend Implementation

- Built in Java 8
 - NDN primitives using jNDN, NDN client library for Java
 - Google's Guice for Dependency Injection
 - Google's Protocol Buffers (Protobuf) for high performance serialization
- Backend is **entirely decoupled** from front end
 - Allows for plugging in of different backend modules
 - Build an IP based backend for direct comparison (future work)

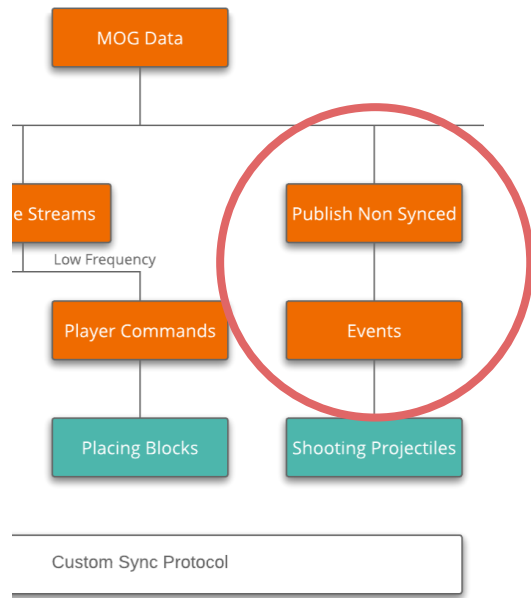
Optimizations: Publisher Side Update Throttling

- Subscribers use extrapolation
 - Publishers can slow down their publish rates
- Publishers use dead reckoning
 - Keep a small cache {nextSN → {PlayerStatus, Timestamp}}
 - Represents data we sent when we specified the next sequence # to be nextSN
- Each subscriber has an outstanding sync interest containing a SN
- Engine publishes update if:
 1. Has no cache entry for nextSN
 2. Local player's velocity has changed*
 3. The **estimate** of the subscriber's position for the local player exceeds a threshold value

* The current implementation's physics engine does not use accelerations meaning changes in velocity are changes in direction 28

Sequence Numbered Cache

- Use case
 - Want to send all events that subscriber has not yet seen
 - **NDN has no information about source of Interest!**
 - But the application has sequence numbers
- Essentially want a multimap like data structure
 - Map of {SN → [new events since that SN]}
 - This is a nice API but would be a very inefficient implementation
- Array based, circular data structure



Testing

Evaluation

Metrics


- Round trip time
- Network traffic
 - Subscriber Interest rates
 - Publisher Interest rates
- PlayerStatus deltas
- Scalability (# nodes)

Parameters

- Caching
 - FreshnessPeriod★
- Topology
- Forwarding strategy
- Publisher throttling
- Area of Interest Management
- Game Mechanics

* Defines hop-by-hop content store lifetime - **not total packet lifetime**


Topologies




2 Game Nodes - Linear



4 Game Nodes - Square



4 Game Nodes - Tree



4 Game Nodes - Dumbbell

Docker

- Basic Docker image containing NFD, NLSR, game and startup script
- Docker **overlay network** to connect all docker nodes together
 - Allows creation of **UDP tunnels** between NFDs running on the nodes
- Deploy locally using docker-compose
- Deploy across cluster using Docker swarms and stacks
- Python scripts
 - Define topologies in code
 - Builds NLSR config file for each node - **this defines the NDN topology**
 - Builds docker compose file for given topology

/NLSR/square/nodeA-nlsr.conf

```
general {
  network /com
  site /stefanolupo
  router /%C1.Router/routerA
}

; Adjacent routers

neighbors {
  neighbor {
    name com/stefanolupo/%C1.Router/routerB
    face-uri udp4://nodeb.ndngame.com
  }
  neighbor {
    name com/stefanolupo/%C1.Router/routerD
    face-uri udp4://noded.ndngame.com
  }
}

; Prefixes we are responsible for

advertising {
  prefix /ndngame/0/discovery/nodeA
  prefix /ndngame/0/nodeA/status
  ;... etc
}
```

topology-builder.py

```
# Build square topology
nodeA = Node("A")
nodeB = Node("B")
nodeC = Node("C")
nodeD = Node("D")

linkNodes(nodeA, nodeB)
linkNodes(nodeB, nodeC)
linkNodes(nodeC, nodeD)
linkNodes(nodeD, nodeA)

topology = "square"
nodes = [nodeA, nodeB, nodeC, nodeD]

NlsrBuilder(topology)
    .buildNlsrFiles(nodes)

ComposeBuilder(topology)
    .buildComposeFile(nodes)
```

docker-compose.yml

```
A:
  image: stefanolupo/ndn:ndn-node
  environment:
    - NLSR_CONFIG=/NLSR/square/nodeA-nlsr.conf
    - GAME=java -jar game.jar -a -hl nodeA
    - METRICS_DIR=/metrics
  hostname: nodeA
  networks:
    private-ndn-overlay:
      aliases:
        - nodea.ndngame.com

# Similar for B, C and D

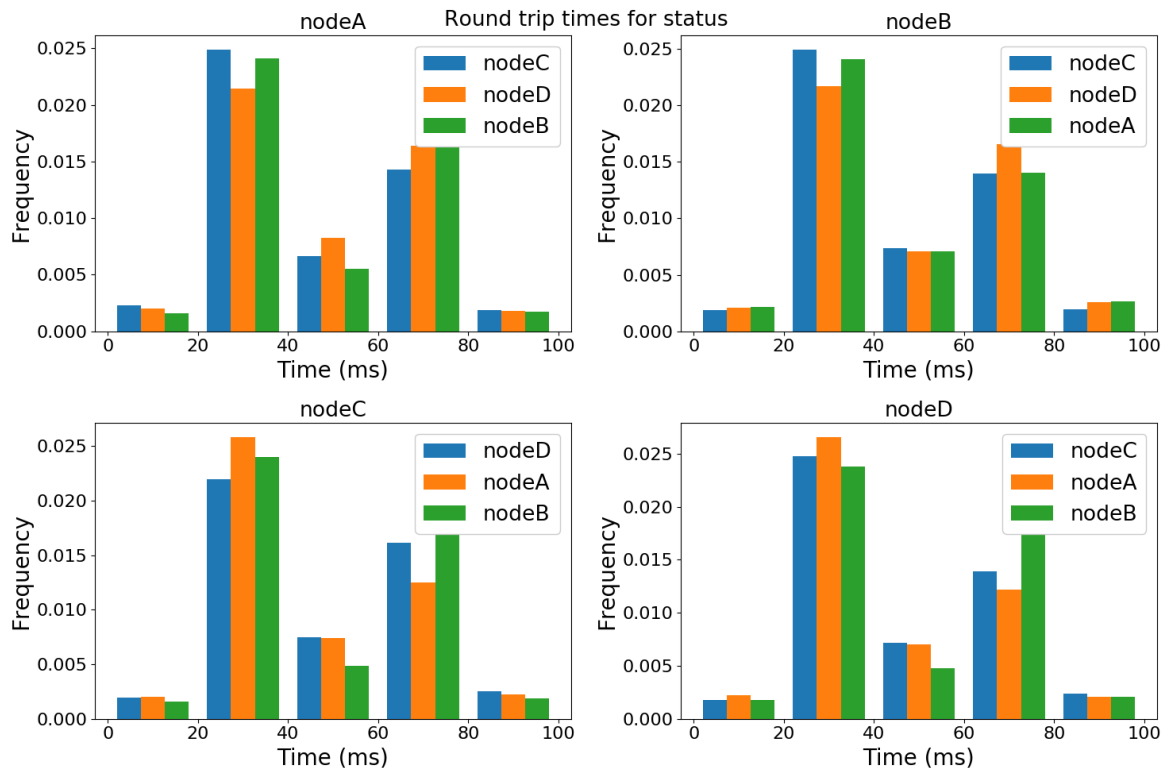
networks:
  private-ndn-overlay:
    driver: overlay
```

Results

Round Trip Times


- ~30ms update rate
- Dumbbell topology
- No caching
- No publish rate tricks

RTT are hugely dependent on publisher update rates




Interest Aggregation

- # Interests **towards** publisher
- # Interests **seen** by publisher



Packets would be required in IP

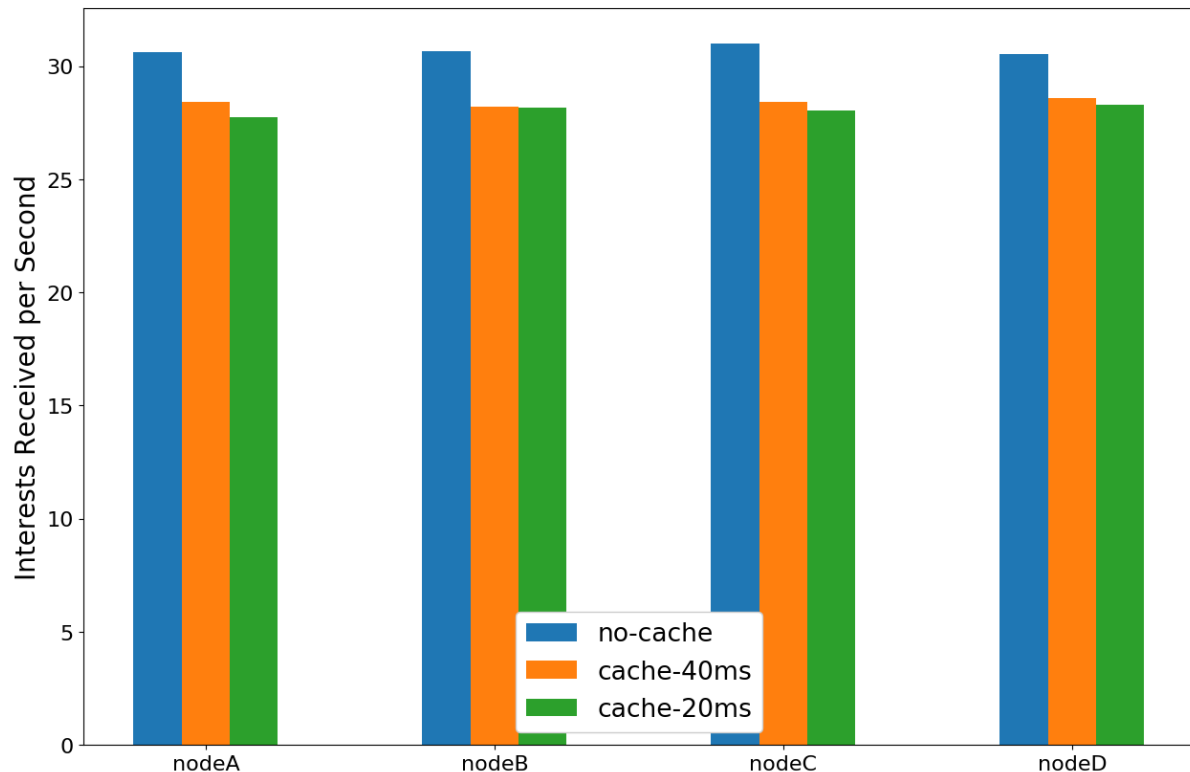
Effects of Interest aggregation on publisher Interest rates



Effects of Caching


- Caching helps during bad periods
- Node can fall slightly behind
- Interest will have already been satisfied
- Receives cached copy

Effects of caching on publisher Interest rates

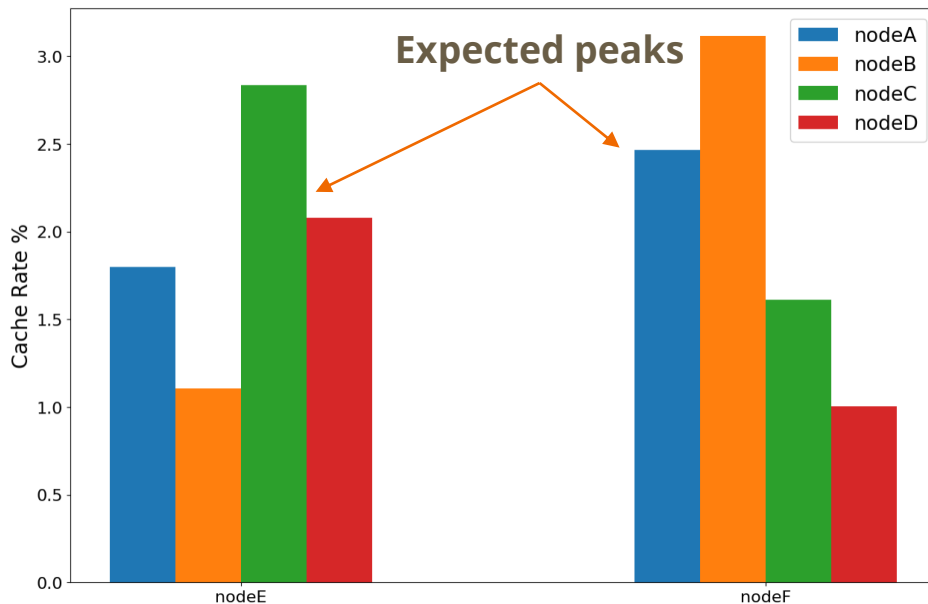


Where caching occurs

- Caching occurs at **intermediate routers** (in dumbbell topology)
- Cache entries **live longer** at intermediate routers
 - Hop by hop **freshnessPeriod**



Router cache rates by node




Publisher Update Throttling

- ~0% cache miss with cache size of 50
- ~45% of updates skippable
- Tolerance of 0.5 GWU



Player Status Updates



Conclusion

Conclusions

- NDN can be very beneficial in P2P settings
 - Interest aggregation and native multicast
 - Caching to a smaller degree
- NDN is still in a proof of concept phase
 - Lots of hidden implementation details making protocol development tricky
- P2P games are much easier to scale
 - However they are very open to cheating

Thank you!